

HSC Chatbot API Documentation

July 26, 2025

Base URL

`http://your-server-ip:8080/`

For local testing:

`http://127.0.0.1:8080/`

Endpoints

GET /

Description:

Renders the chat HTML page. This is the frontend interface for interacting with the chatbot.

Response:

Returns an HTML page.

POST /get

Description:

Accepts user input from a form and returns a response generated via a Retrieval-Augmented Generation (RAG) pipeline.

Request Parameters:

- `msg` (string): The user's question/input query.

Example Request (form-data or x-www-form-urlencoded):

```
POST /get
msg=Who is nelson mandela?
```

Backend Workflow

1. PDF Loader:

- Loads and splits content using `RecursiveCharacterTextSplitter`
- Source: `/Data` folder

2. Embeddings:

- Model: `sentence-transformers/all-MiniLM-L6-v2`

3. Vector Store:

- Qdrant Cloud (`hsc_bot_qdrant2`)
- Similarity search (`k=5`)

4. Language Model:

- Qwen/Qwen3.0.6B via Hugging Face Transformers

HSC Chatbot API Documentation (Continued)

- Integrated with LangChain using HuggingFacePipeline

Setup & Run

To start the server:

```
python app.py
```

Ensure:

- PDF files exist in /Data/
- Qdrant (if used) is loaded properly
- The model and dependencies are installed
- Port 8080 is open (if testing remotely)

Deprecation Warnings

Update imports:

```
from langchain_community.vectorstores import Qdrant
from langchain_community.embeddings import HuggingFaceEmbeddings
from langchain_community.document_loaders import PyPDFLoader,
    DirectoryLoader
```

Security Notes

- Do not expose your Qdrant API key in public code or URLs.
- Use environment variables and .env files.
- The current server runs in development mode - not safe for production.